

Operating Systems: IPC & Synchronization

Fabrice BOISSIER <fabrice.boissier@epita.fr>



2021-11-09



How to manipulate data in IT

Two visions of the world:

- You know in advance the length of your data
 - Record length, block size, ...
 - Specific OS and languages (*z/OS and JCL, COBOL, ...*)
 - Specific functions (*strnlen, strncpy, strncmp, ...*)
 - Pretty easy to manage (and somewhat safe)
- You know the separator of each field/record/file/...
 - « , », « ; », « \n », « \0 », « EOF », ... (*CSV format, awk(1), ...*)
 - Hard to manage: your CPU manipulates N bits, buffers of fixed length, ...
 - Way more unsafe...
 - Foundations of modern computing

IPC & Synchronization

- Why divising the work between multiple processes?
- Functional reasons:
 - Multiple organizations are responsible of a precise service....
...and a product is using the services of those organizations
- Technical reasons:
 - Multiple machines with one main service on each one
(We must link them and make them communicate/transfer data)
 - One machine composed of multiple processors or cores of execution

IPC & Synchronization

- How to make processes communicate and share or transfer data from one to another?
 - InterProcess Communications
- What are the technical challenges?
 - Who will read which data at what time?
 - How to wait for the other process to finish reading/writing?
 - How to be sure that the data are written immediately?
 - ...
 - *Synchronization problems*

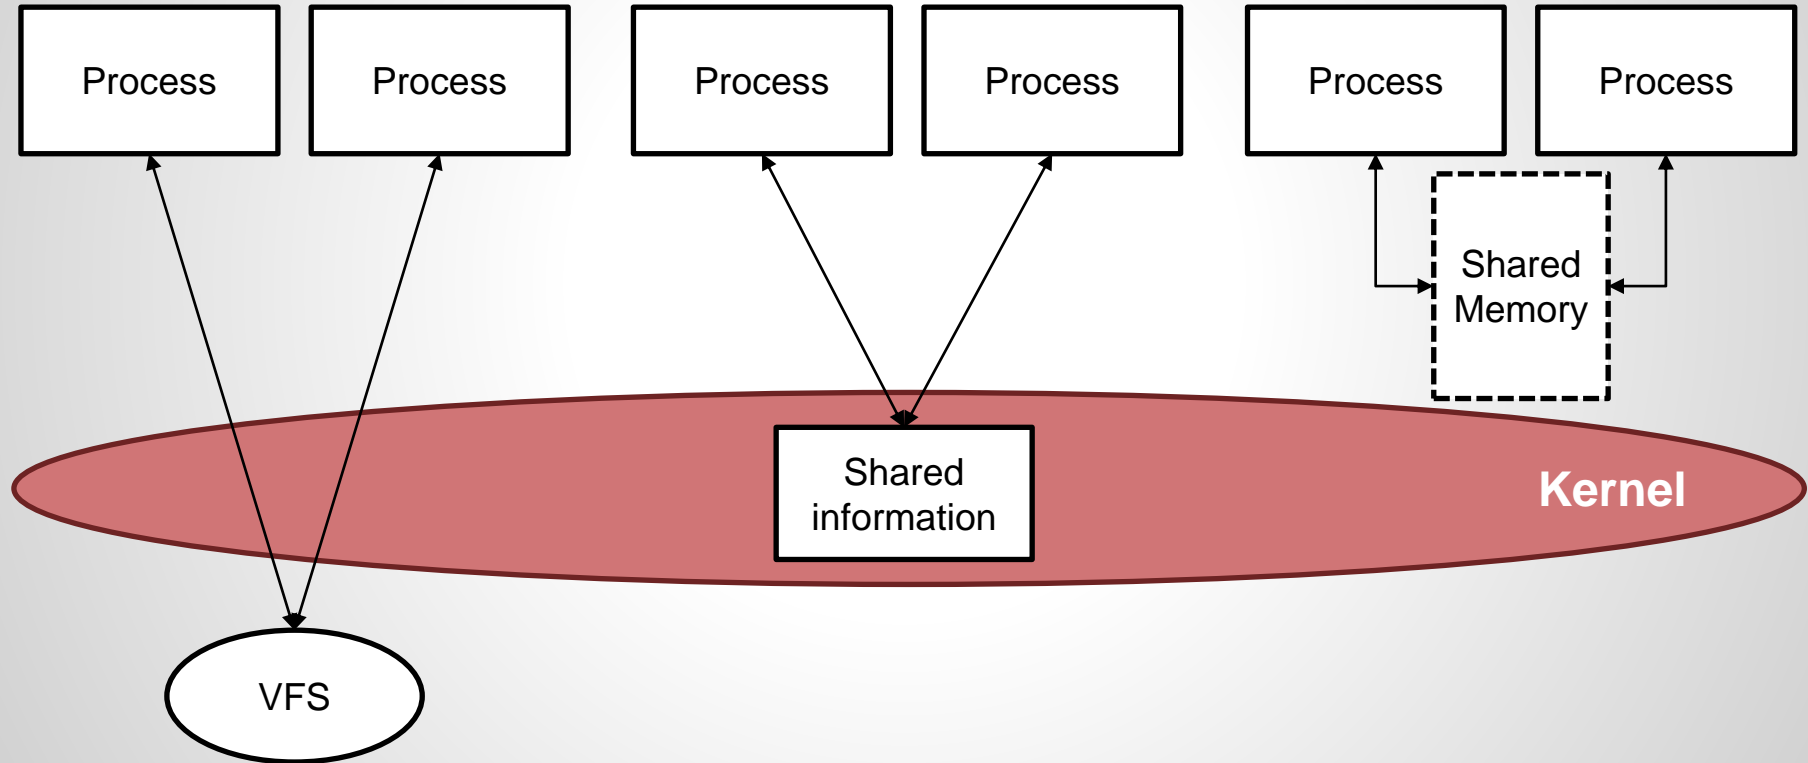
IPC & Synchronization

UNIX Network Programming, Volume 1: The Sockets
Networking API - *W. Richard Stevens*

UNIX Network Programming, Volume 2: Interprocess
Communications - *W. Richard Stevens*

InterProcess Communication

InterProcess Communication



InterProcess Communication

- File
 - Processes read and write within a file
 - Requires to pass through the VFS and in the kernel (if monolithic)
- Shared informations
 - Pipe, Message queue, Semaphore
 - Processes ask the kernel to share data and manage the accesses to it
- Shared memory
 - Shared memory
 - Memory pages are shared among processes
 - Kernel is invoked « only » when creating/destroying the pages

IPC: Problems & Issues

- Problems

- Reading and writing data concurrently might corrupt data
(synchronization in read/write is required)
- Passing through the kernel is expensive
(Reduce the number of syscalls)

- Issues

- How to establish a link (rendez-vous)
- How many processes per link (1 to 1, 1 to N)
- How many links per process

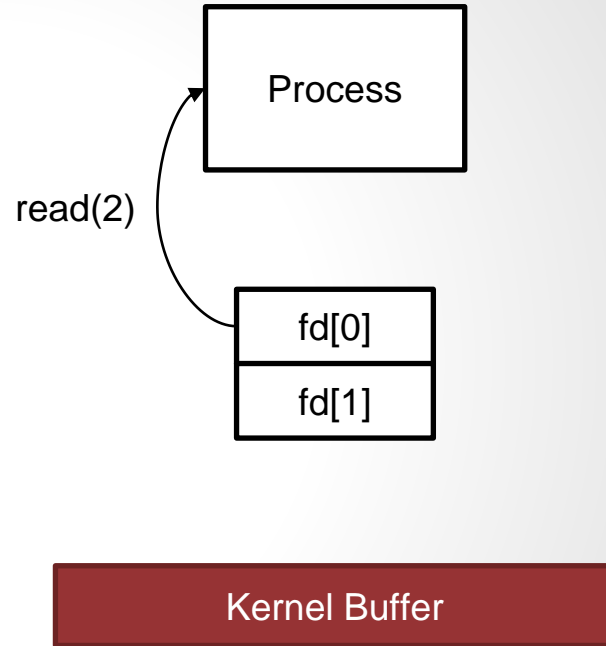
IPC: Pipes (anonymous pipes)

- A buffer is created within the kernel
 - Works exactly like a FIFO
 - Bytes are read in the some order they were written
- Access with read/write through a file descriptor
 - A VFS « file » is created and points to the kernel buffer
 - « Everything is a file »
- Each read « consumes » data
 - Data is not present anymore within the pipe after a read
- *Used with a fork(2)*

IPC: Pipes (anonymous pipes)

```
int pipe(int fd[2])
```


- Read on fd[0]



IPC: Pipes (anonymous pipes)

```
int pipe(int fd[2])
```

- Read on fd[0]
- Write on fd[1]




Kernel Buffer


IPC: Pipes (anonymous pipes)

```
int pipe(int fd[2])
```

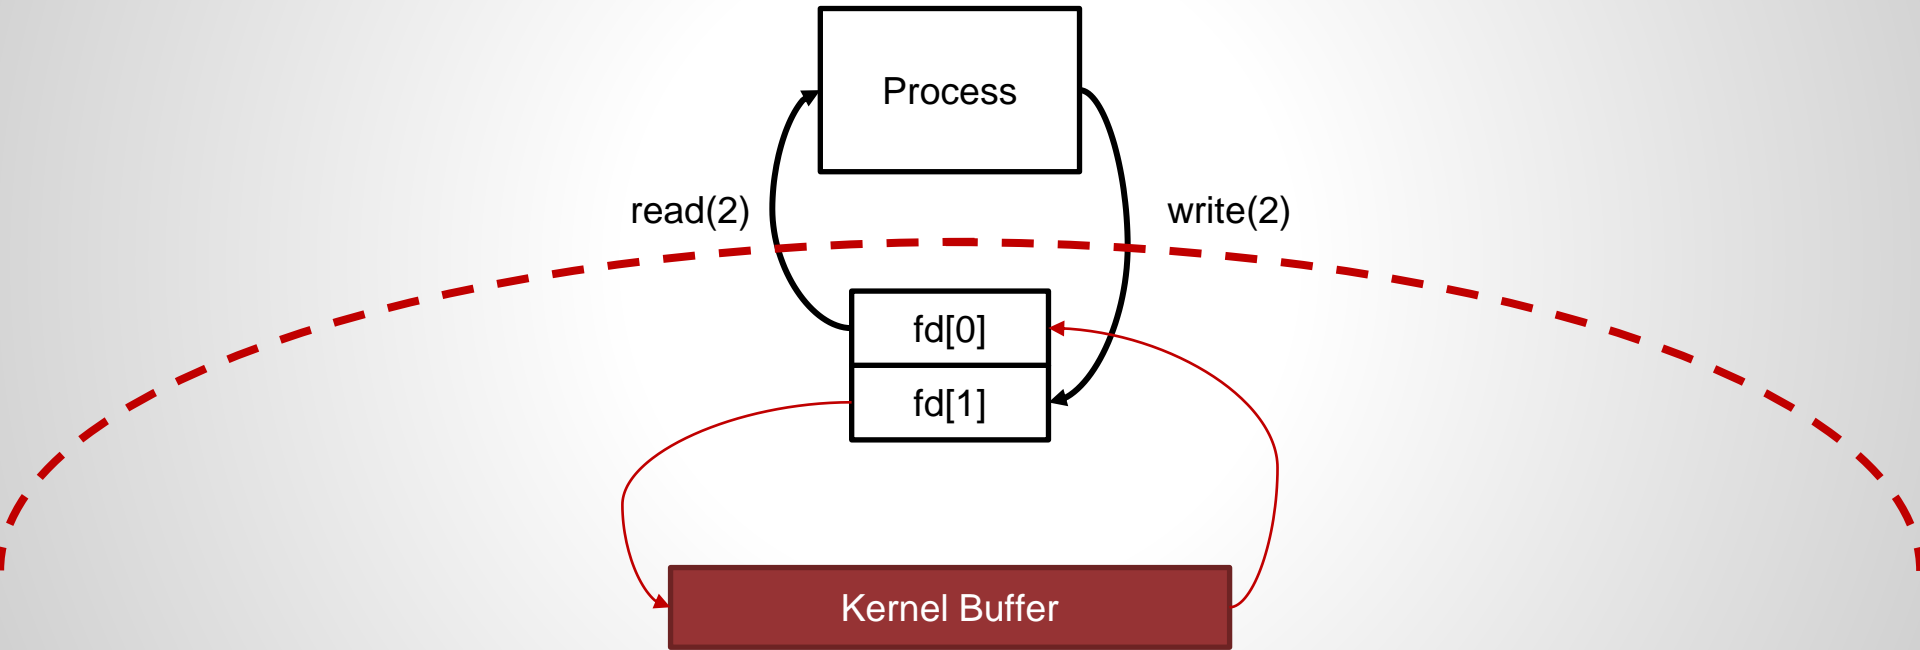
- Read on fd[0]
- Write on fd[1]
- The kernel manages the transfer of data



IPC: Pipes (anonymous pipes)




IPC: Pipes (anonymous pipes)




IPC: Pipes (anonymous pipes)


fork(2)



IPC: Pipes (anonymous pipes)



IPC: Pipes (anonymous pipes)




IPC: Pipes (anonymous pipes)

- Kernel manages synchronization
 - *Except if the read/write is asynchronous...*
- When there is no data available, the process that read is put in a waiting state
 - The scheduler is waiting for the resource to be available before putting back the process within the *ready queue*
- When the buffer is full, the process that writes is put in a waiting state


IPC: Pipes (anonymous pipes)

You should close the useless file descriptors BEFORE writing or reading!



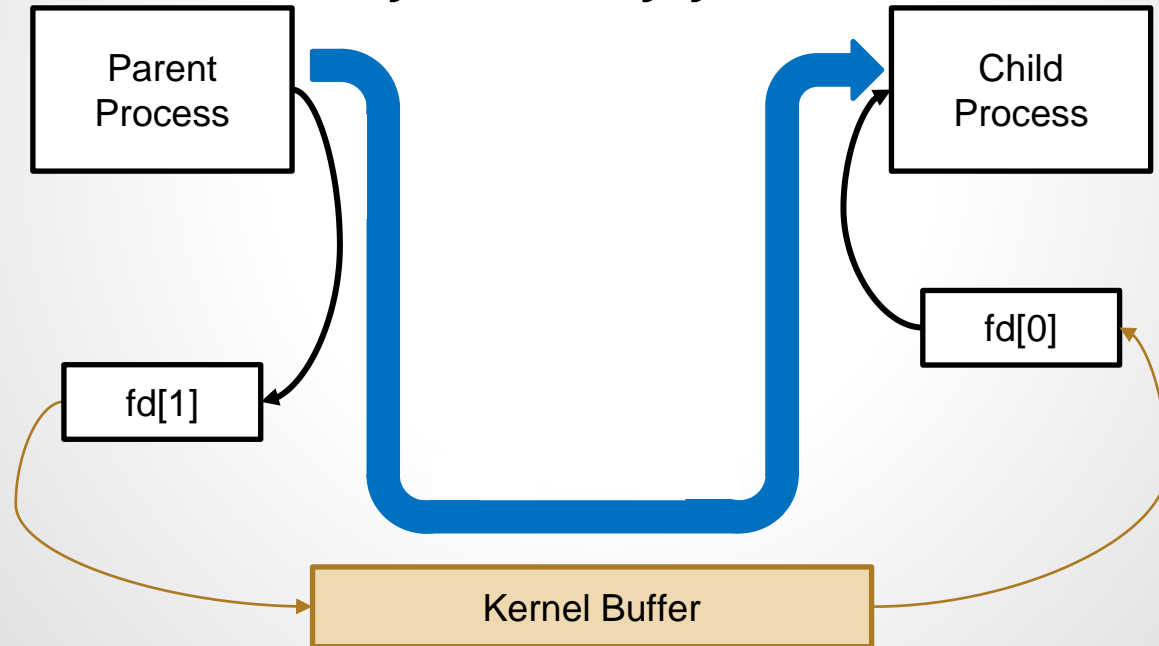
IPC: Pipes (anonymous pipes)

In the pipe context, you should keep a very clean state of your file descriptors



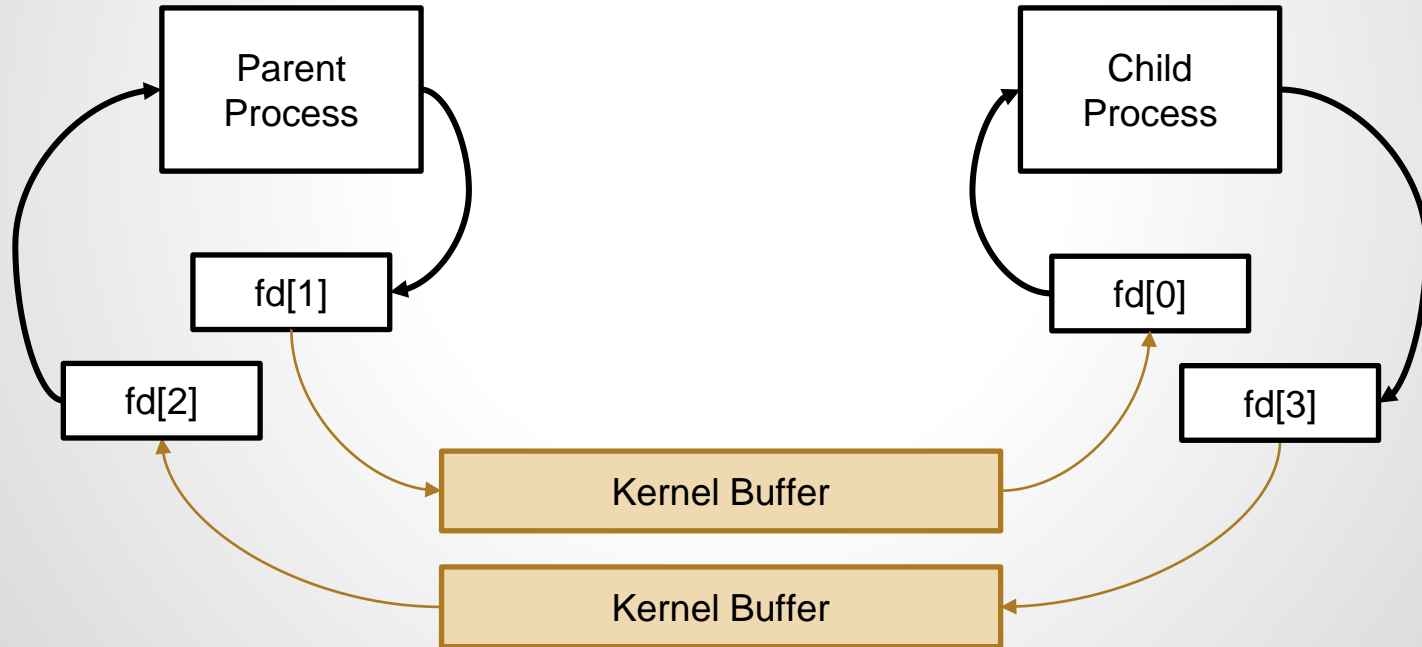
IPC: Pipes (anonymous pipes)

*Anonymous pipes are traditionally « unidirectional »
(read very carefully your manual)*



IPC: Pipes (anonymous pipes)

*For bidirectionnal, you must create 2 pipes
(just create a table for 4 int to put the fd)*



IPC: Pipes (anonymous pipes)

- Kernel buffer size:
 - Currently 16 memory pages (64KiB or 65,536 Bytes)...
...it might be 1 Memory Page (4KiB / 4,096B) in some cases...
...or it can be dynamically chosen
(*check with « ulimit -a » or fcntl(2) with F_GETPIPE_SZ*)
 - Pipe buffer size can be updated with *fcntl(2)* and *F_SETPIPE_SZ*
 - Writing more than the buffer size will block the writing process
- Pipes are linked to the VFS
 - Each fd points to a VFS « file » for keeping read/write cursor state
 - « pipefs » on Linux
(*specific management of the buffers*)

IPC: Pipes (anonymous pipes)

- Well, now you know how to use anonymous pipes...
 - See `pipe(2)` for details
- ...but how is implemented the « | » in the shell ?
 - See `sh(1)` for « some » details...
 - ...see `strace(1)` for much more details

[syscalls & strace: let's find what's happening]

- `strace(1)`
- Traces syscalls (even within childs)
 - Might check for specific syscalls
- Useful for understanding « how » things work...
- ...and why they do not work in some cases...
- ...or even how to break things.

IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```

```
strace -qf -e execve,pipe,dup2,read,write \
  sh -c 'cat file.txt | grep "hello" '
```

- Let's search for these syscalls
 - « execve », « pipe », « dup2 », « read », « write »
- « -f » even within the childs

```
pipe([3, 4])  
[... fork() fork() ...]  
[pid 27597] dup2(3, 0)  
[pid 27597] execve("/bin/grep", ["grep", "hello"], ...  
[pid 27596] dup2(4, 1)  
[pid 27597] read(3, "", 4096)  
[pid 27596] execve("/bin/cat", ["cat", "file.txt"], ...  
[pid 27596] read(3, "", 4096)  
[pid 27597] read(0, <unfinished ...>  
[pid 27596] read(3, "hello world!\n", 131072)  
[pid 27596] write(1, "hello world!\n", 13)  
[pid 27597] <... read resumed> "hello world!\n", 32768)  
[pid 27597] write(1, "hello world!\n", 13)  
[pid 27597] read(0, <unfinished ...>  
[pid 27596] read(3, "", 131072)  
[pid 27597] <... read resumed> "", 32768)  
[pid 27597] +++ exited with 0 +++  
[pid 27596] +++ exited with 0 +++
```

IPC: Dup2

(<https://toroid.org/unix-pipe-implementation>)


- `dup2(int oldfd, int newfd)`
 - See *dup(2)* and *dup2(2)*
- Duplicates a file descriptor into another one
 - The VFS « file » pointed in the fd table is exactly the same
 - Therefore, they share the exact same state (lseek, flags, ...)
- « oldfd » is duplicated into « newfd »
 - If « newfd » was opened, it is first closed
 - If « oldfd » is not valid, nothing else is done, and `dup2()` fails
 - If « oldfd » == « newfd », nothing happens

IPC: Dup2

(<https://toroid.org/unix-pipe-implementation>)

- `dup2(int oldfd, int newfd)`
 - See *dup(2)* and *dup2(2)*
- Main usage of `dup2()`: replace one or more file descriptor entry from `STDIN`, `STDOUT`, or `STDERR` to another VFS file before an `exec(2)`
- Don't forget to close the « `oldfd` » after the duplication
 - If you do an `exec(2)`, it won't close the « `oldfd` »...
 - ...neither will the next program (except if it is coded to do so)

IPC: Dup2




Step 1: *(initial state)*

- The 3 usual file descriptors are opened

- STDIN
- STDOUT
- STDERR

(They are inherited from the parent process)

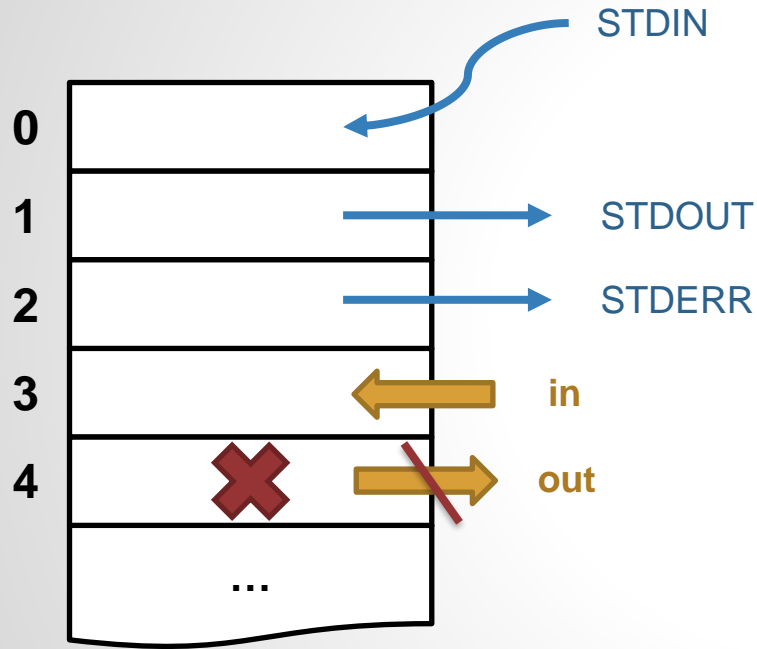
IPC: Dup2



Step 2: `pipe(int[2])`

- Two file descriptors are reserved in order to access the pipe (the kernel's buffer)


IPC: Dup2



Step 3: `close(4)`

- One of the file descriptor is closed, in order to keep the pipe unidirectionnal


IPC: Dup2



Step 4: `dup(3, 0)`

- Duplicate the file descriptor number 3 into STDIN
1. fd 0 is first closed
 2. fd 3 is redirected to the fd 0

IPC: Dup2




Step 5: `close(3)`

- Close the old file descriptor (fd 3) because the program that will be loaded with `exec()` won't close it

(by inheritance, each process expects only fd 0,1,2 to be opened)

IPC: Dup2



Step 6: *(final state)*

- The next program to be executed will have three opened file descriptors (0 in reading mode, 1 and 2 in writing mode), but the file descriptor 0 is connected to a pipe instead of STDIN

```
pipe([3, 4])  
[... fork() fork() ...]  
[pid 27597] dup2(3, 0)  
[pid 27597] execve("/bin/grep", ["grep", "hello"], ...  
[pid 27596] dup2(4, 1)  
[pid 27597] read(3, "", 4096)  
[pid 27596] execve("/bin/cat", ["cat", "file.txt"], ...  
[pid 27596] read(3, "", 4096)  
[pid 27597] read(0, <unfinished ...>  
[pid 27596] read(3, "hello world!\n", 131072)  
[pid 27596] write(1, "hello world!\n", 13)  
[pid 27597] <... read resumed> "hello world!\n", 32768)  
[pid 27597] write(1, "hello world!\n", 13)  
[pid 27597] read(0, <unfinished ...>  
[pid 27596] read(3, "", 131072)  
[pid 27597] <... read resumed> "", 32768)  
[pid 27597] +++ exited with 0 +++  
[pid 27596] +++ exited with 0 +++
```

IPC: Pipes (anonymous pipes) and dup2

```
pipe([3, 4])
```

- pipe(2) created file descriptors « 3 » and « 4 »

```
[pid 27597] dup2(3, 0)
```


- dup2(2) duplicated fd « 3 » into fd « 0 » (STDIN)

```
[pid 27596] dup2(4, 1)
```

- dup2(2) duplicated fd « 4 » into fd « 1 » (STDOUT)


IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

sh\$> cat file.txt | grep "hello"




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



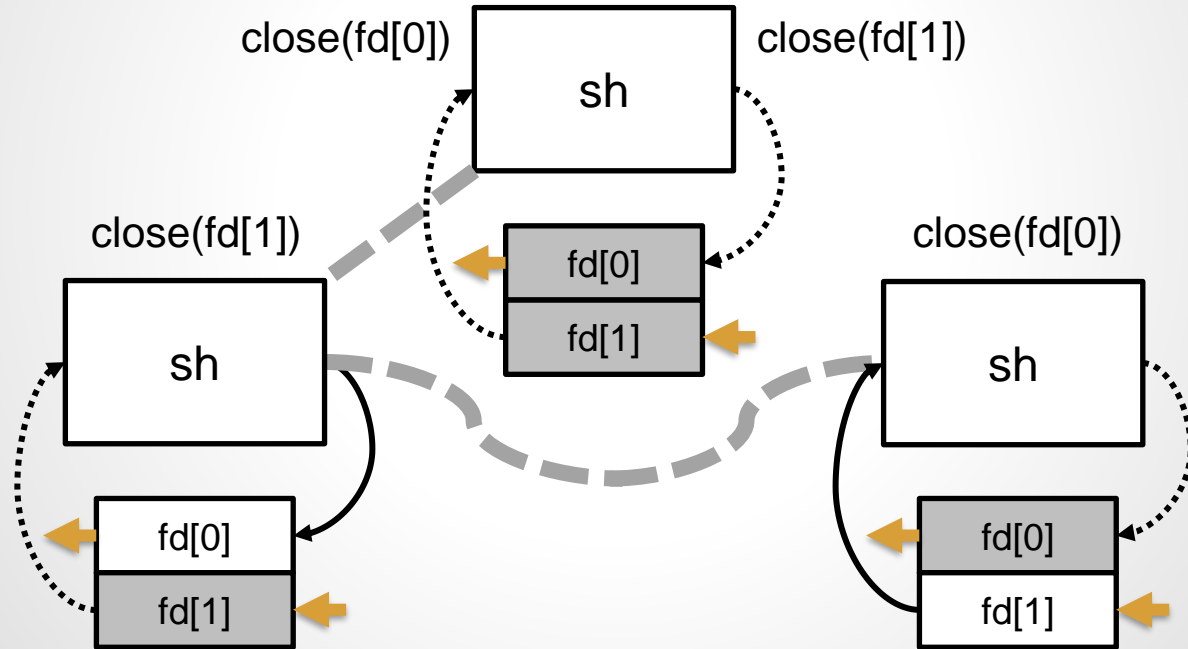
IPC: Pipes in the shell

sh\$> cat file.txt | grep "hello"




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



IPC: Pipes in the shell

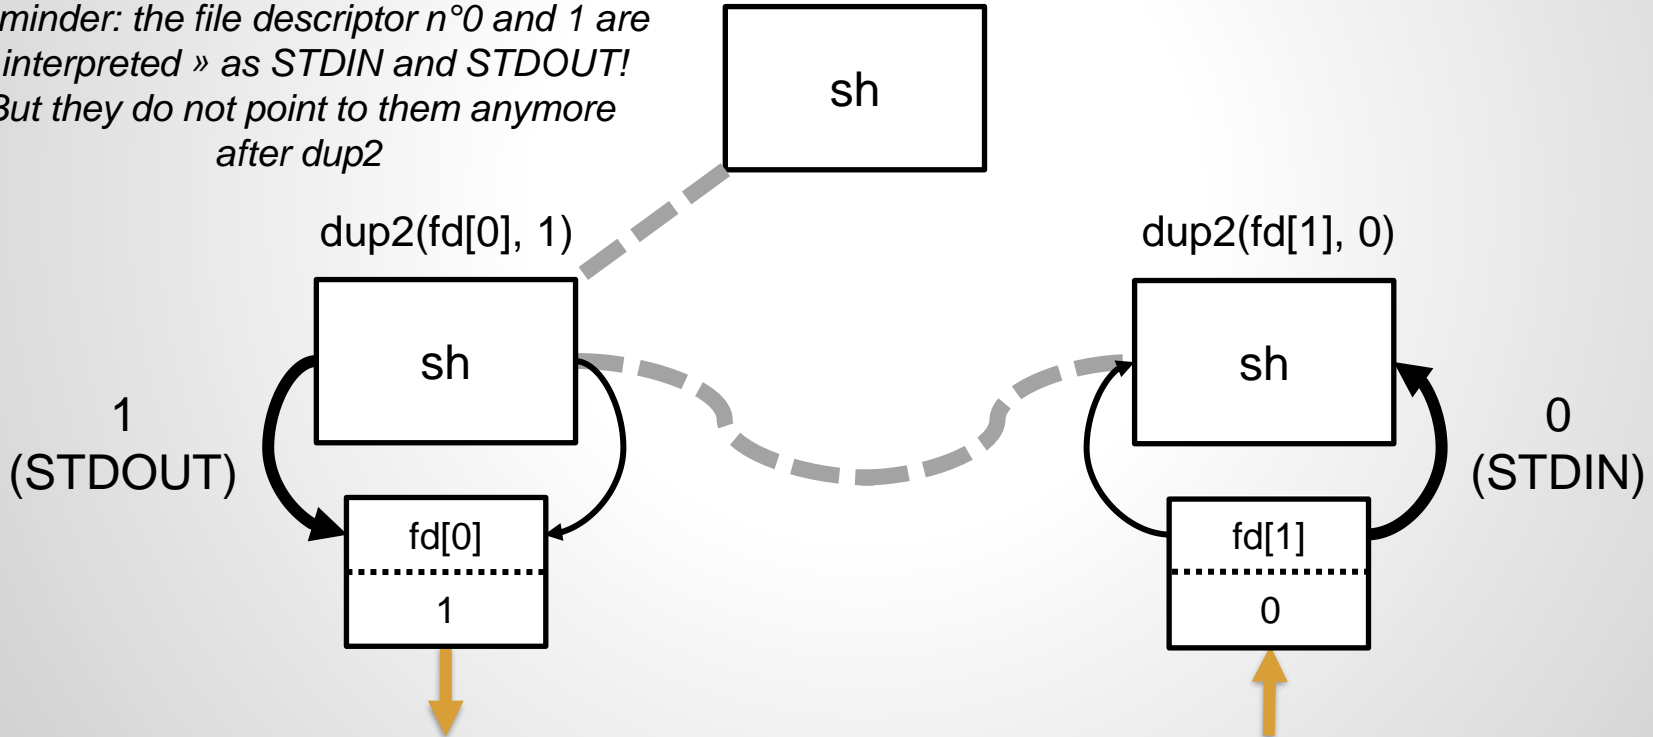
```
sh$> cat file.txt | grep "hello"
```



IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```

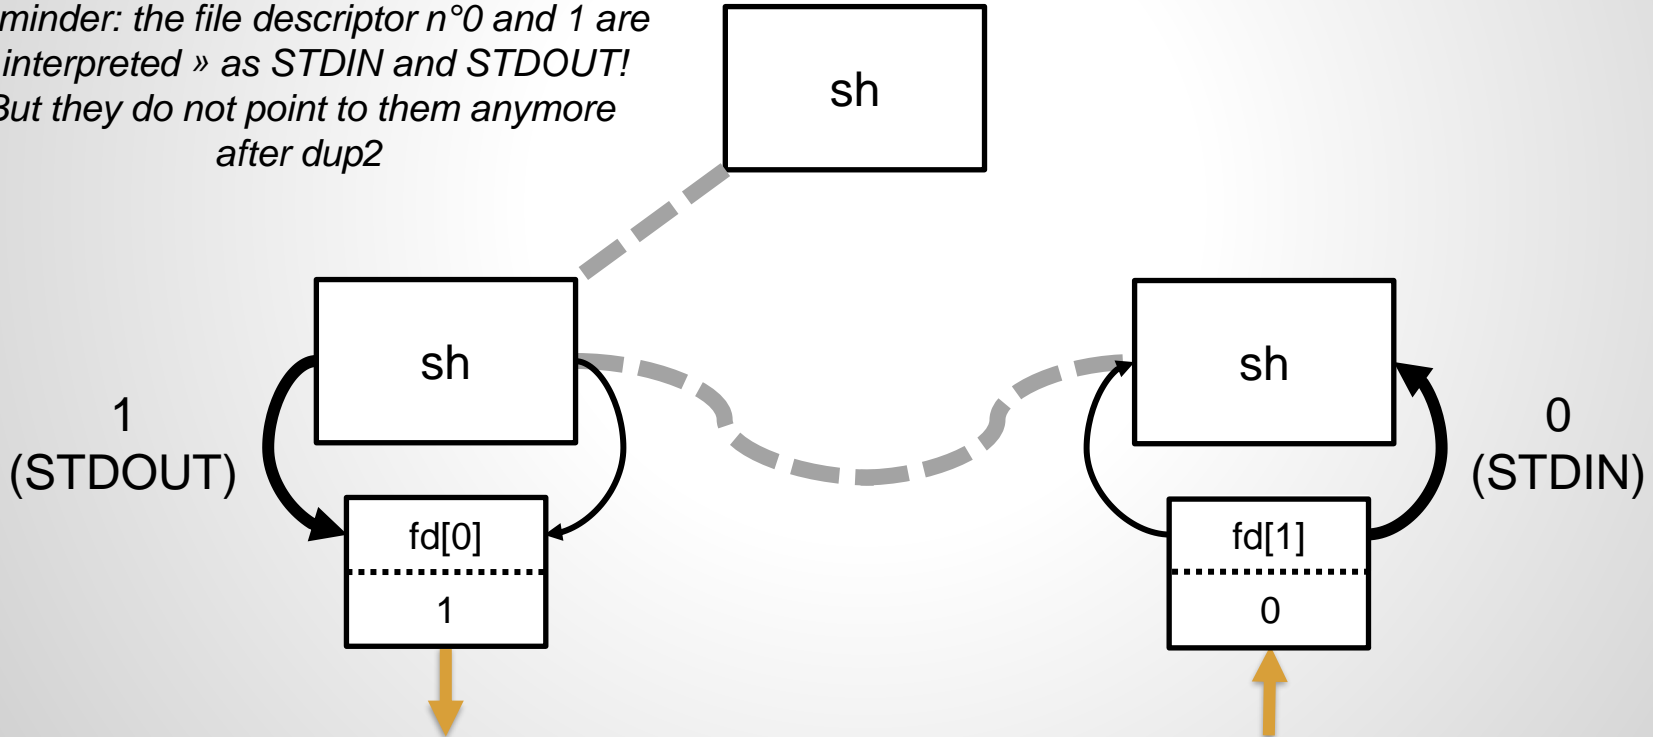
*Reminder: the file descriptor n°0 and 1 are
« interpreted » as STDIN and STDOUT!
But they do not point to them anymore
after dup2*



IPC: Pipes in the shell


```
sh$> cat file.txt | grep "hello"
```

*Reminder: the file descriptor n°0 and 1 are
« interpreted » as STDIN and STDOUT!
But they do not point to them anymore
after dup2*




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



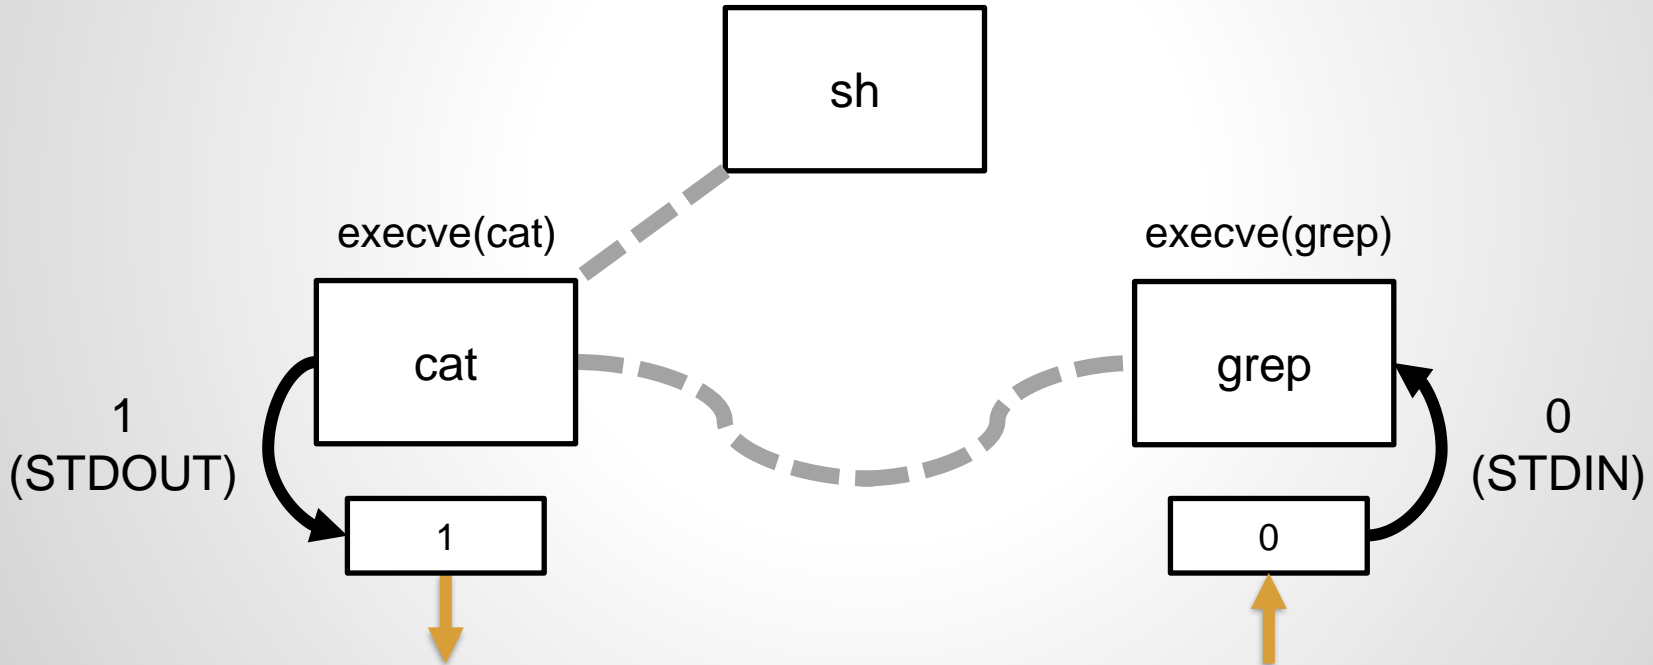
IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



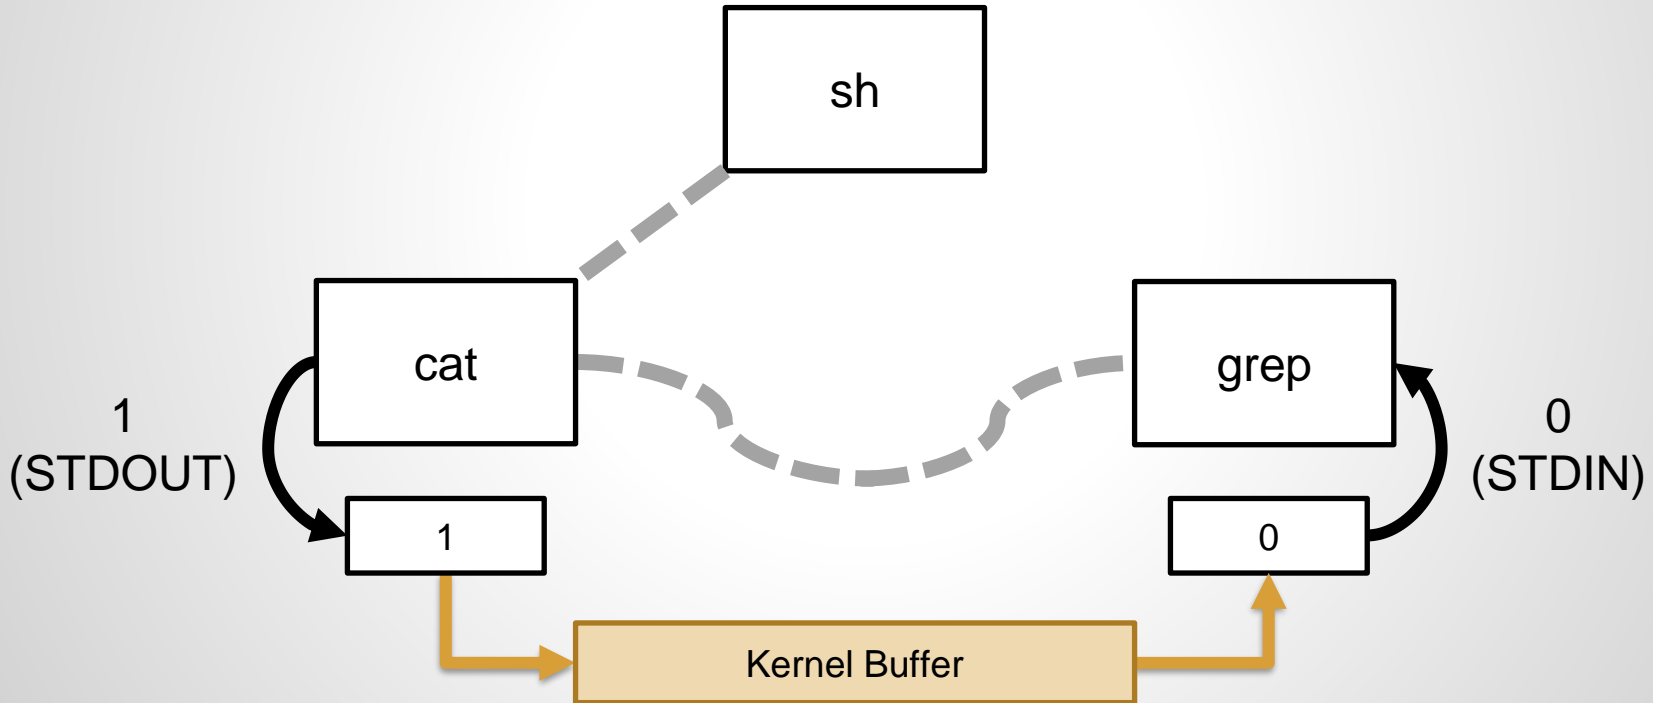
IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



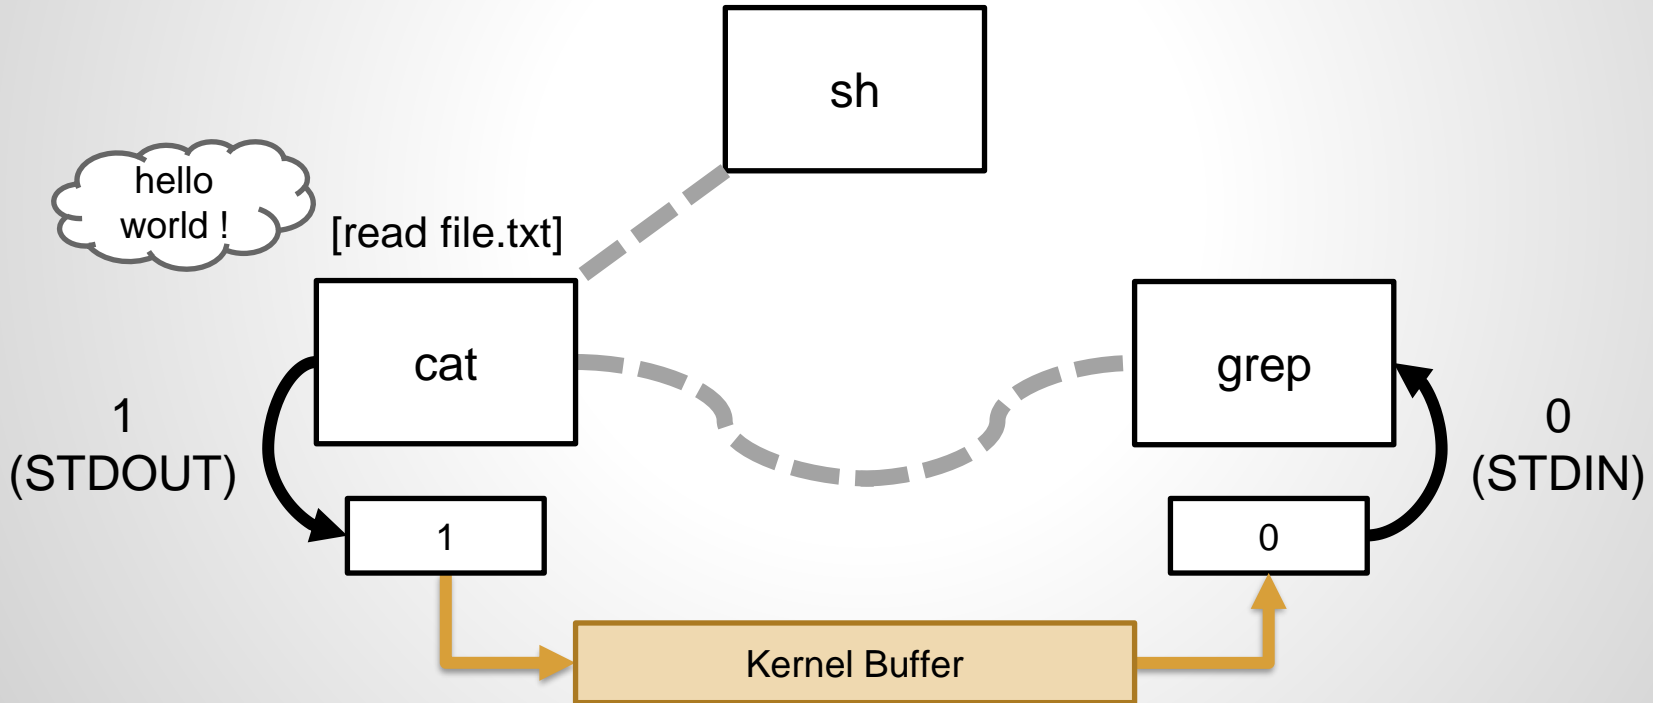
IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



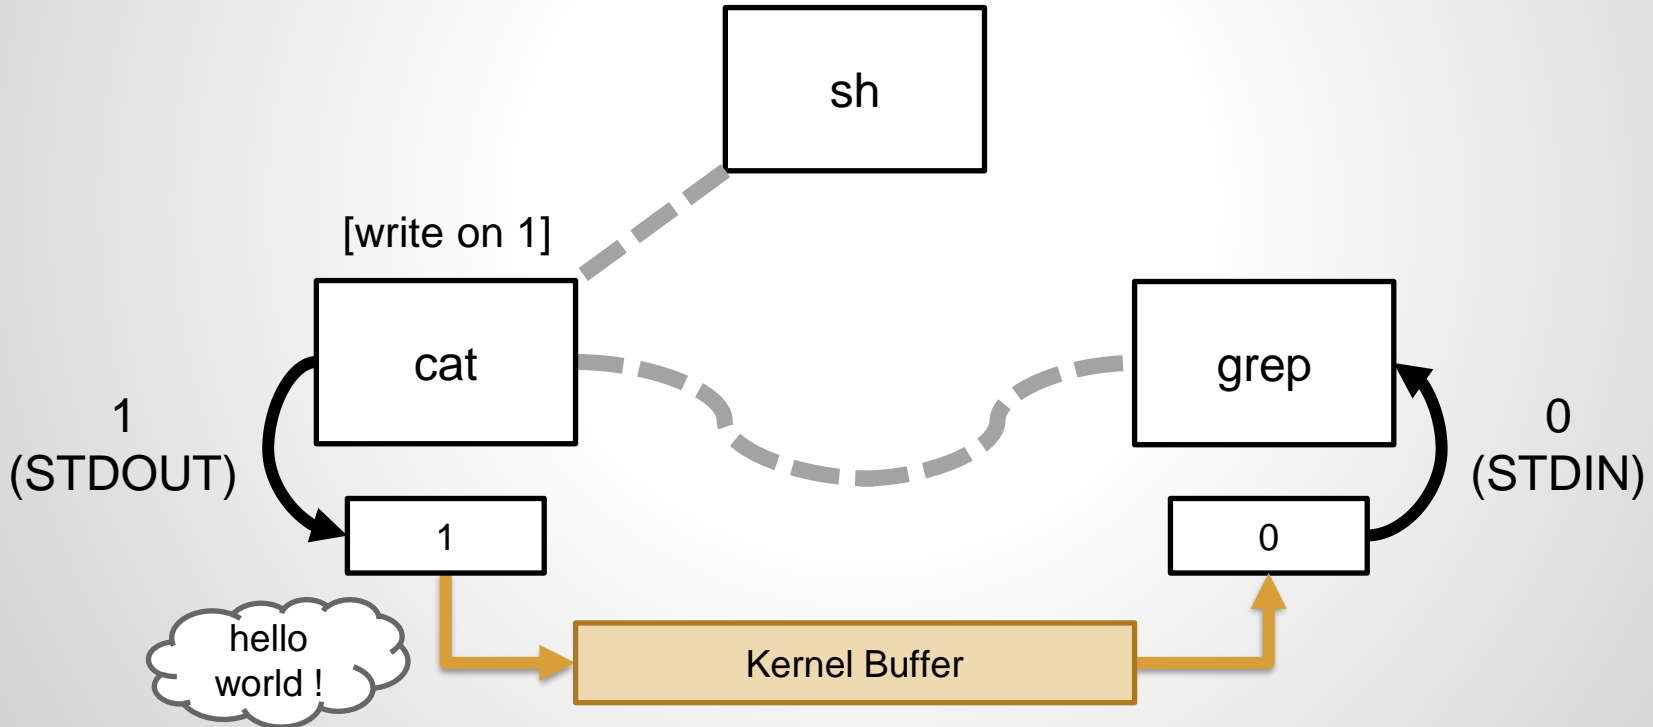
IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



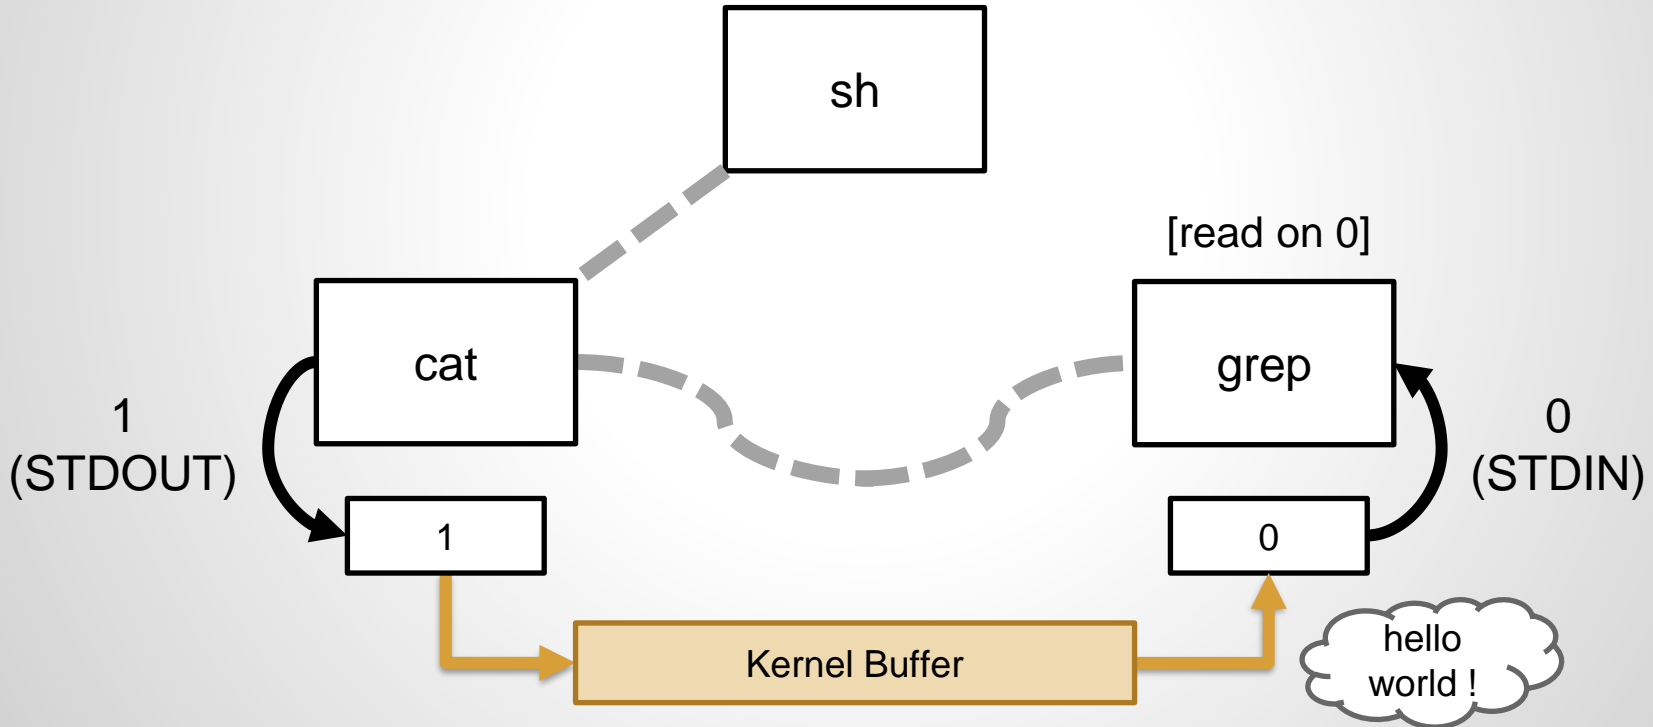
IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



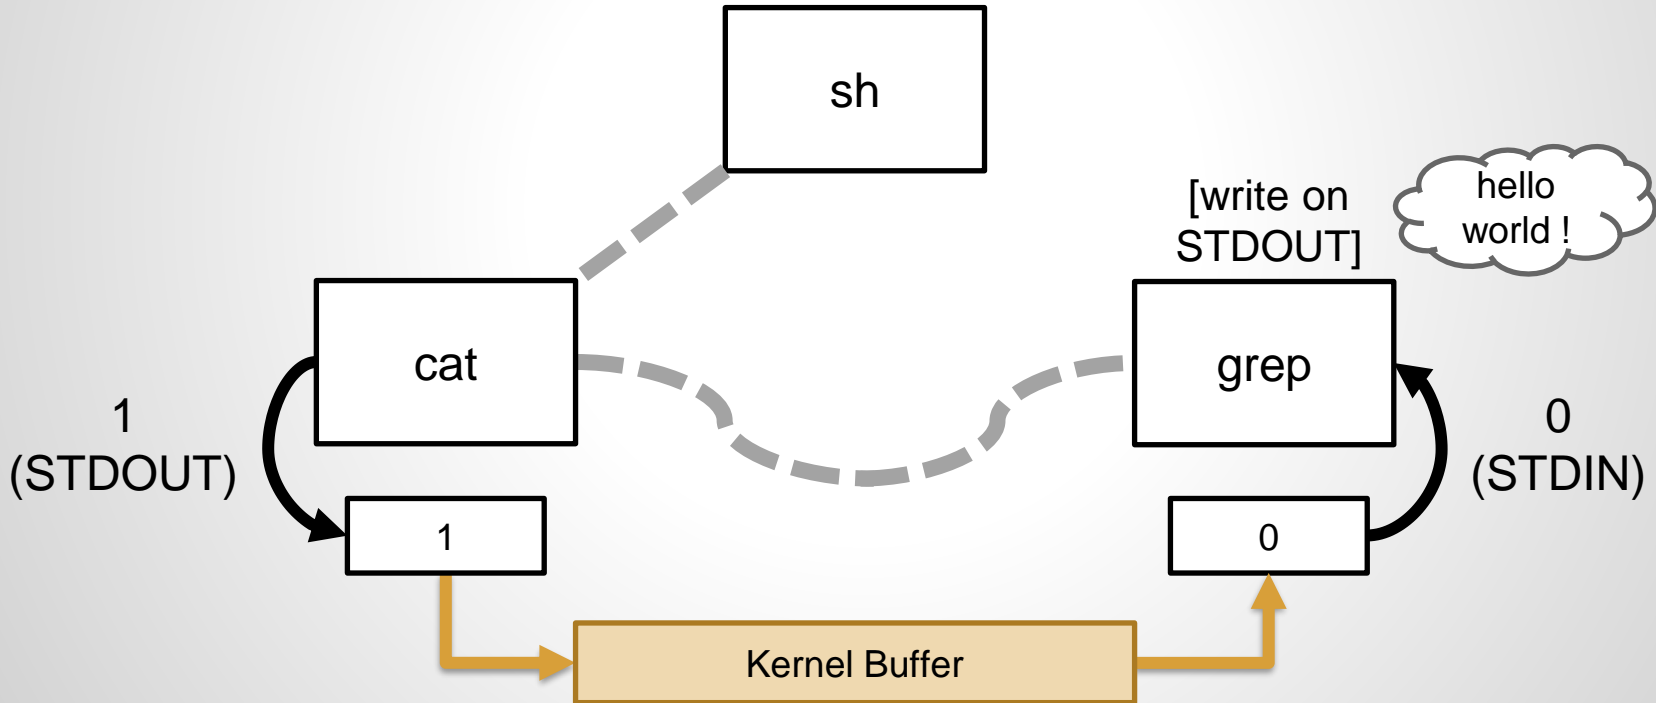
IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



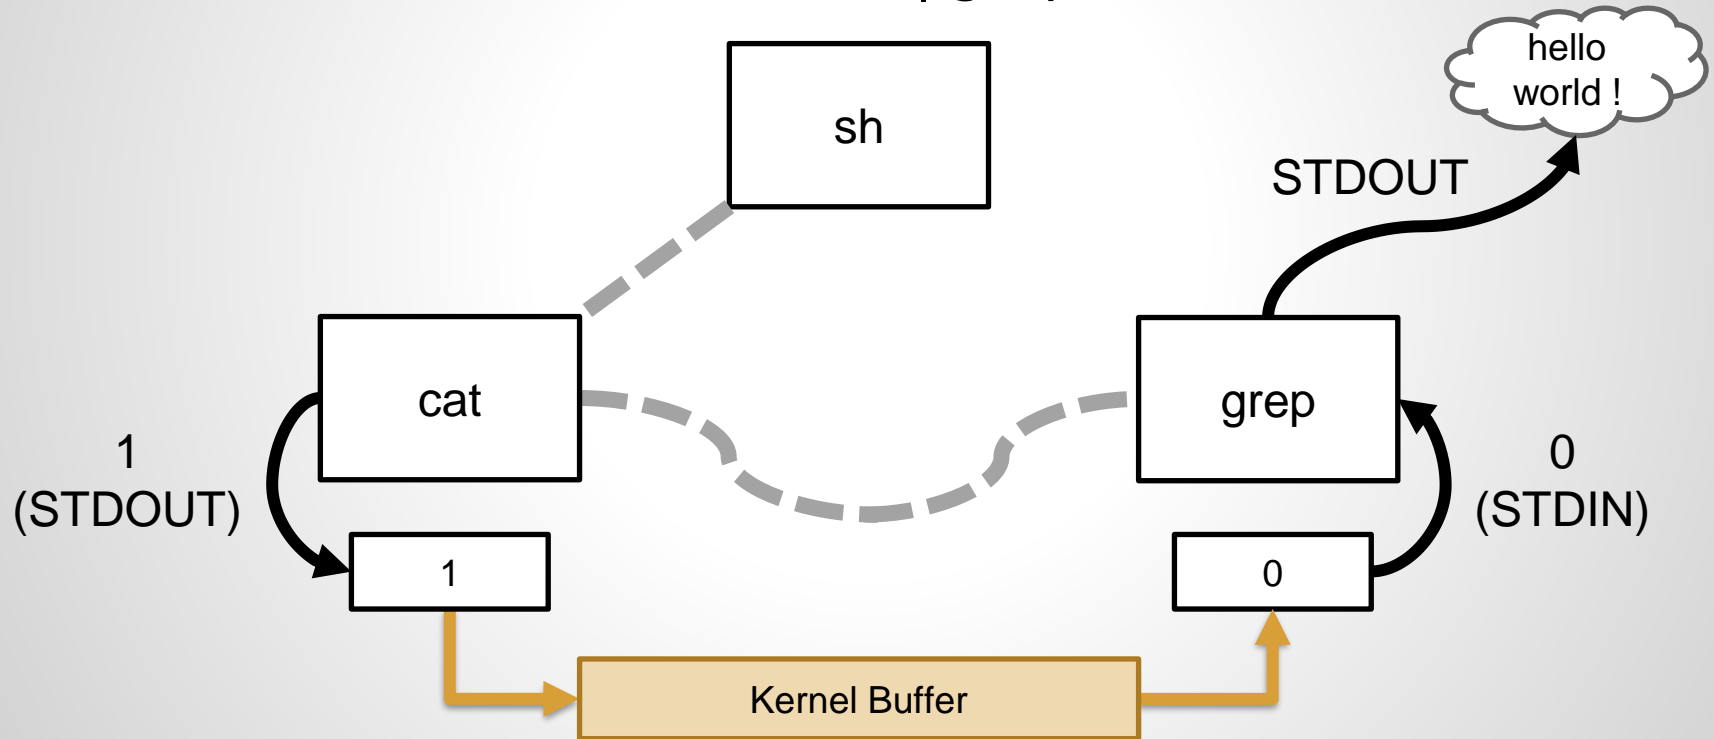
IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```




IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



IPC: Pipes in the shell

```
sh$> cat file.txt | grep "hello"
```



IPC: Named pipes

- Anonymous pipes are created between processes that have a close common parent (or directly between a parent and its child)
- Named pipes are specific files stored on the file system
 - Processes search for a specific pathname in the file tree
 - *See `mkfifo(1)` and `mknod(2)`*
- Named pipes can be used for communication between unrelated processes (no close common parent)
 - But each process must know the name of the named pipe

IPC

- Other IPCs require a clear rendez-vous point
 - `ftok(2)`: A « key » is used as the name of the rendez-vous point
 - `ipcs(1)`: prints some of the IPCs currently in use
 - `ipcrm(1)`: removes some active IPCs
- Messages Queues (q)
 - Bi-directionnal communication with messages (type & data)
- Shared Memory (m)
 - Memory shared among processes
- Semaphores (s)
 - Used for synchronization

IPC: Messages Queues

- Mailbox of messages
- A message is a structure with a type, and data
 - Messages have a precise length
- Creation / Destruction: *msgget(2)* / *msgctl(2)*
 - Defines a clear name for the messages queue (the « key »)
 - Might be restricted to child processes only
- Send / Receive: *msgsnd(2)* / *msgrcv(2)*
 - Reception is blocking if no message is available
 - The message structure 1st field & msgrcv 4th parameter allow to choose which message get (*see in the manual how to use them smartly*)

IPC: Shared Memories

- Memory shared among multiple processes
- Synchronization is *required*
 - See semaphores and mutexes...
- Creation / Destruction: *shmget(2)* / *shmctl(2)*
 - Defines a clear name for the shared memory (the « key »)
 - Decides the length of the memory part to share
- Attach / Detach: *shmat(2)* / *shmdt(2)*
 - Uses a pointer to the shared memory
 - *fork(2)* shares the memory between parent and child

IPC: Semaphores

(*synchronization*)

- Counter shared among multiple processes
 - Cannot be below 0
 - Must use specific routines to change the value
 - *sem_overview(7): named and unnamed semaphores*
- Creation / Destruction: *sem_init(3)* *sem_open(3)* / *sem_destroy(3)*
 - Defines a clear name for the shared memory (the « key »)
 - Defines the maximum value of the counter
- Increase / Decrease: *sem_post(3)* / *sem_wait(3)*
 - *sem_post(3)*: increases the counter
 - *sem_wait(3)*: decrements the counter OR waits until it can be decreased


Quick overview of the threads

Multithreading

- Problems: How to...
 - Allows parallelism inside a process?
 - Reduces the cost of context switching?
- Solution
 - Thread (lightweight process): state, registers & stack
Shares other resources
 - Process: group of threads.
Classical process = process with only 1 thread
- Functionalities
 - Same as a process: creation, termination, state, etc...
 - New issues: concurrent access on shared resources

```
1. char *myStr;  
2. int i = 0;  
3. int main(void)  
4. {  
5.     const char *var = "Test";  
6.     int a = 1337;  
7.     i = addition(21, 42);  
8.     myStr = malloc(32 * sizeof(char));  
9.     return (0);  
10. }
```



Process



```

1. char *myStr;
2. int i = 0;
3. int main(void)
4. {
5.     const char *var = "Test";
6.     int a = 1337;
7.     i = addition(21, 42);
8.     myStr = malloc(32 * sizeof(char));
9.     return (0);
10. }

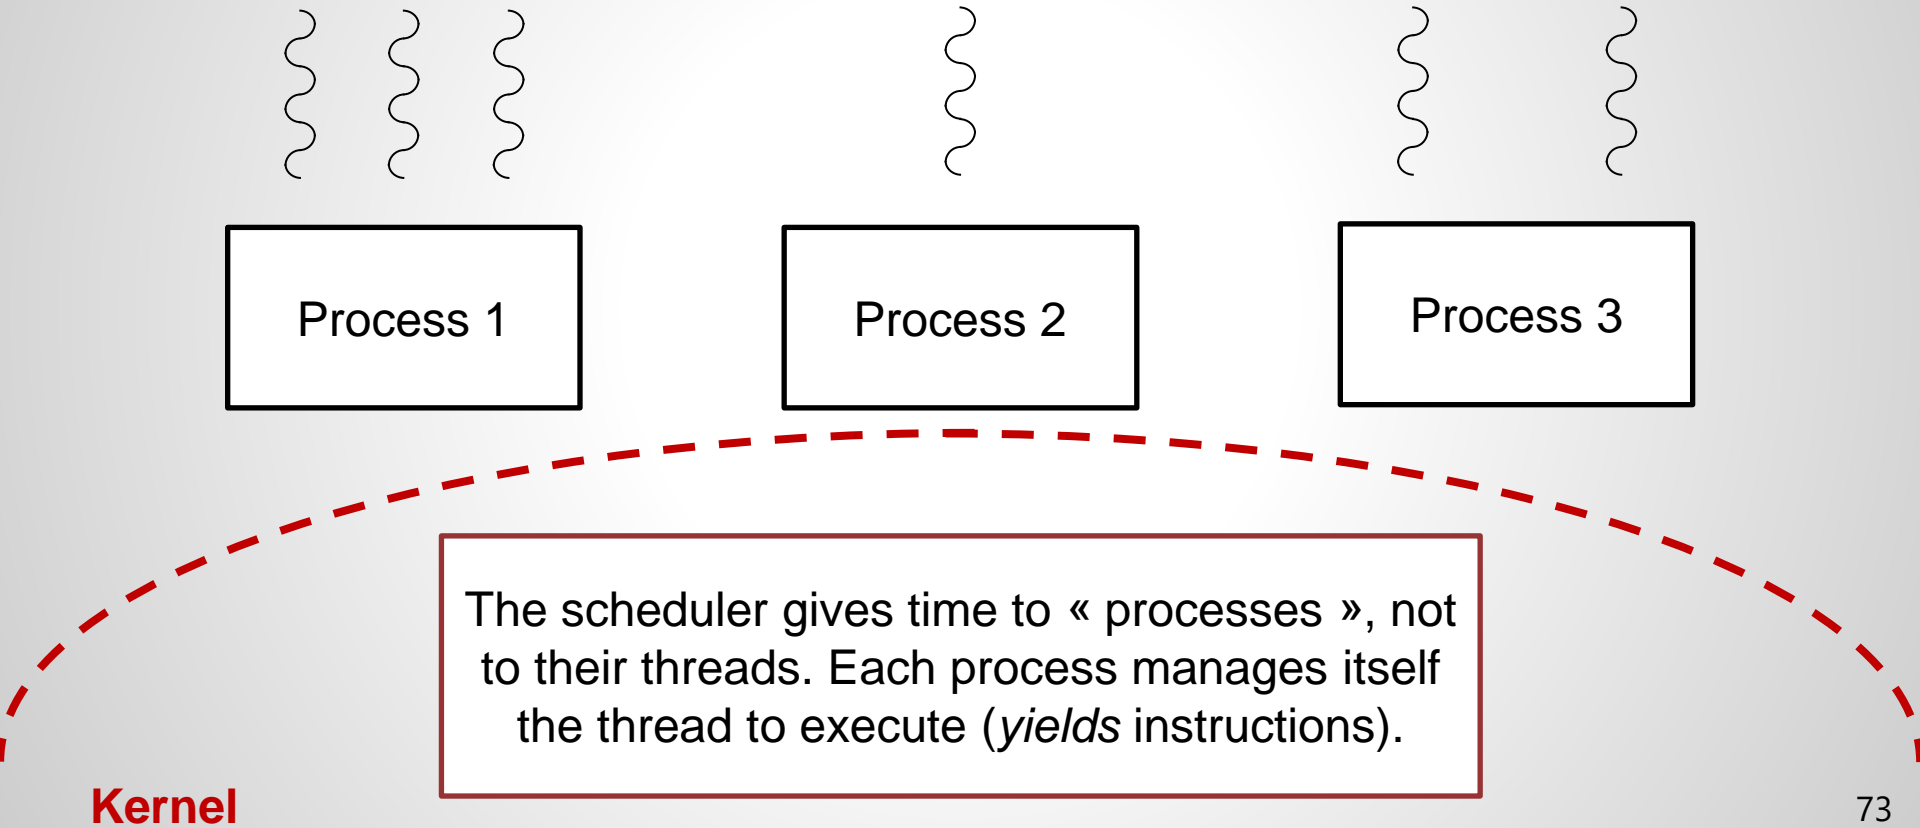
```



Userland Threads

- Principle
 - Implemented as a library in userland
 - 1 thread table per process (managed by the library)
- Pros
 - Usable on a system without support for threads
 - Fast context switching (no kernel trap)
 - Customizable scheduling algorithm
- Cons
 - Needs for unblocking syscalls
 - Threads can lock the CPU (they need to yield explicitly)
 - Threads are used to alleviate blocking

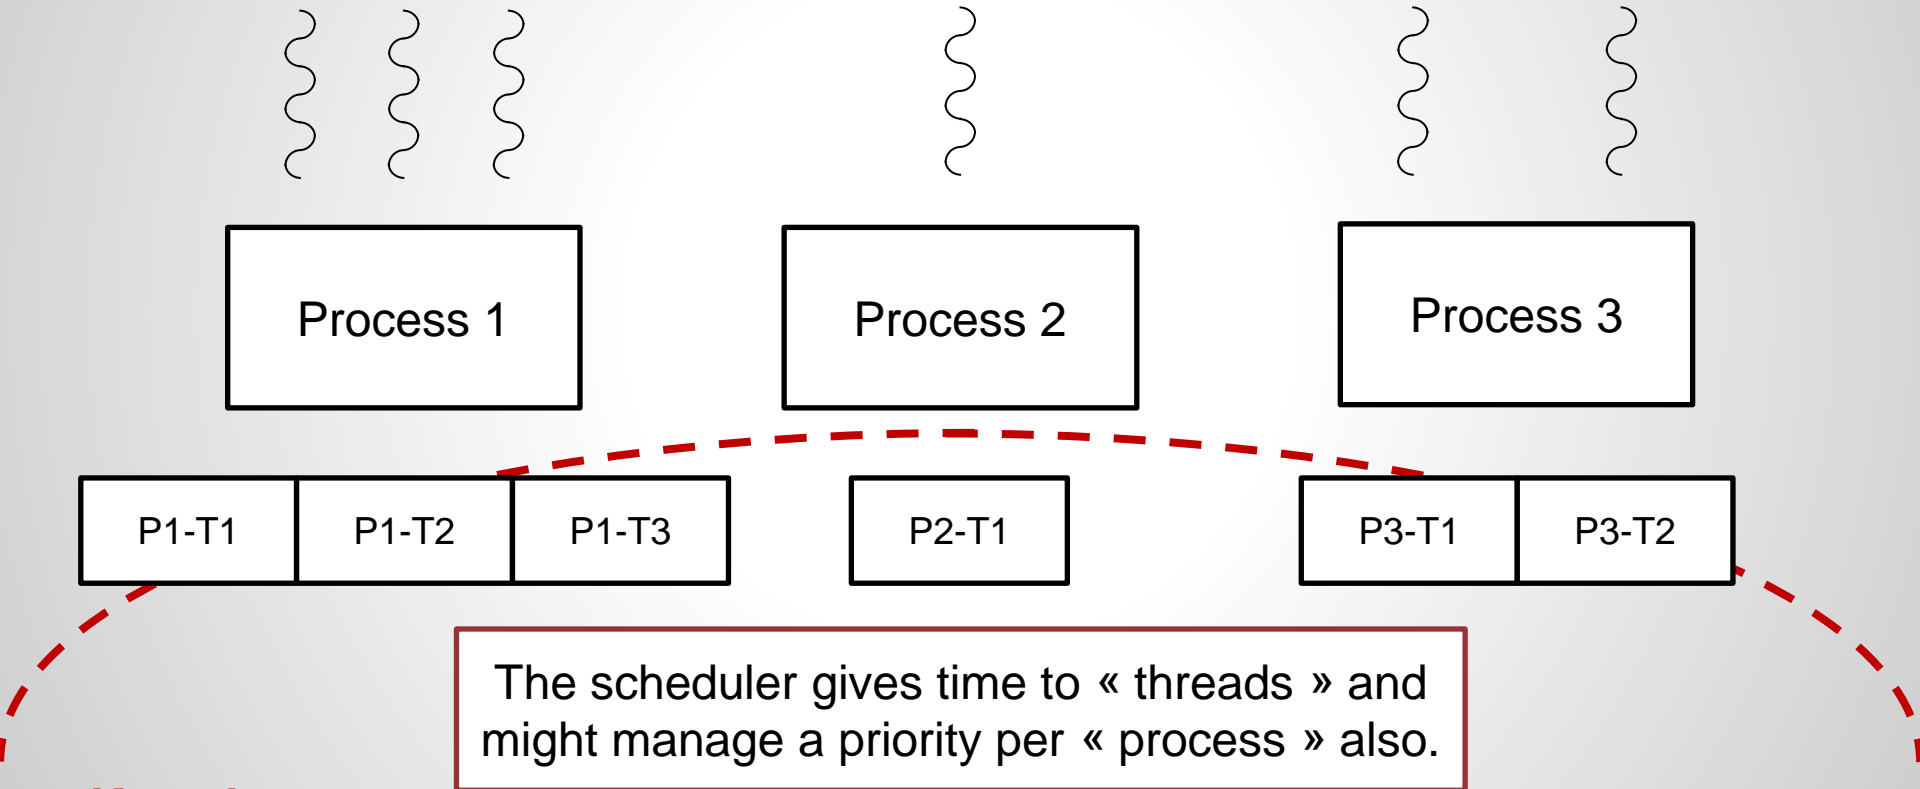
Userland Threads



Kernel Threads

- Principle
 - Adds a thread table inside the process table (used by the scheduler)
 - Every blocking call is implemented as a syscall
- Pros
 - Ease to create an application using them
 - No need for non blocking calls
- Cons
 - Creation/deletion/bookkeeping have a cost
 - Interrupt & blocking syscalls

Kernel Threads



Unified API: Pthread

- POSIX API used to run threads
- Simple unified interface for multi threaded environment on POSIX system
- Beware: everything is shared between threads...
 - Except the thread ID in the scheduler => each thread is independent
- pthreads(7)
 - pthread_create(3), pthread_join(3), pthread_yield(3), ...

Where? What?

Per Thread

- Thread ID
- Signal mask
- Errno
- Scheduling policy
- Capabilities
- CPU affinity

Per Process

- Process ID
- Parent Process ID
- Process Group
- User/Group ID
- File descriptors
- umask
- Current directory
- Limits
- ...

Concurrency problems

Concurrency

*A ressource is shared among multiple processes...
...they all want to access it simultaneously!*

*How to manage this case?...
...well, is there only one case of concurrency?*

Concurrency

- One resource to share
 - Available / Used
- How many processes or threads try to access it?
 - Maximum of users
- Usual processing:
 1. Access to the resource (lock)
 2. Process
 3. Release of the resource (unlock)

Concurrency

- Mutual exclusion (mutex)
- Ressource allocation
- Producers - Consummers
- Readers - Writers

Mutual exclusion

```
Booking()  
{  
    if (AvailableSeats > 0)  
    {  
        /* Processing for booking */  
        AvailableSeats -= 1;  
    }  
}
```

- If 1 process executes this code, *AvailableSeats* will stay at 0 or above
- If 2 processes executes simultaneously this code...
...In some cases, *AvailableSeats* might fall below 0

Mutual exclusion: critical section

```
Booking()  
{  
  
    if (AvailableSeats > 0)  
    {  
  
        /* Processing for booking */  
  
        AvailableSeats -= 1;  
  
    }  
  
}
```

- This part is a *critical section*
 - only « one » process (or thread) can execute it if you wish to stay in a coherent state
- 3 properties:
 - Mutual exclusion
 - Bounded wait
 - Progress

Mutual exclusion: critical section

- Mutual exclusion
 - Only one process/thread can be « inside » the critical section
- Bounded wait
 - A process/thread waiting to access the critical section must be able to execute it within a maximum waiting time
- Progress
 - If no process/thread is « inside » the critical section, nothing must block its access, and the processes/threads waiting to enter must decide which one will enter first

Mutual exclusion: Blocking interrupts

```
Booking()
{
    disable_interrupt();

    if (AvailableSeats > 0)
    {
        /* Processing for booking */
        AvailableSeats -= 1;
    }

    enable_interrupt();
}
```

- Blocking interrupts
 - Useful for kernel processing like scheduling or cache protection
- Pros:
 - Cannot interrupt the execution while in the critical section
- Cons:
 - The process must be supervisor
 - All the processes become unaware of interruptions...
 - Do not work on multi core/CPU₈₅

Mutual exclusion: Test and Set

```
int *lock = 0;
```

```
Booking()
```

```
{
```

```
    while (TAS(lock) == 1)
```

```
    { }
```

```
    if (AvailableSeats > 0)
```

```
    {
```

```
        /* Processing for booking */
```

```
        AvailableSeats -= 1;
```

```
    }
```

```
    (*lock) = 0;
```

```
}
```

- Test and Set (*TAS*)
 - Atomic operation
- Pros:
 - The resource is really protected
- Cons:
 - Busy waiting for others
 - Do not ensure fair access to critical section

Mutual exclusion: Test and Set

```
int TAS(int *lock)
{
    disable_interrupt();

    int test;

    test = (*lock);
    (*lock) = 1;
    return (test);

    enable_interrupt();
}
```

```
// A high level vision of TAS
// (it must be "atomic")
```

```
int *lock = 0;

Booking()
{
    while (TAS(lock) == 1)
    { }
    if (AvailableSeats > 0)
    {
        /* Processing for booking */
        AvailableSeats -= 1;
    }
    (*lock) = 0;
}
```

Mutual exclusion: Compare and Swap

```
int *lock = 0;
```

```
Booking()  
{
```

```
    key = 1;
```

```
    while (key == 1)
```

```
        SWAP(lock, key);
```

```
    if (AvailableSeats > 0)
```

```
    {
```

```
        /* Processing for booking */
```

```
        AvailableSeats -= 1;
```

```
    }
```

```
    (*lock) = 0;
```

```
}
```

- Compare and Swap (CAS)
 - Swaps two values
 - Atomic operation
- Pros:
 - The resource is really protected
- Cons:
 - Busy waiting for others
 - As all the accesses are in a readable memory

Mutual exclusion: Swap example

```
int SWAP(int *A, int *B)
{
    disable_interrupt();

    int tmp;

    tmp = (*B);
    (*B) = (*A);
    (*A) = tmp;

    enable_interrupt();
}
```

// A high level vision of CAS
// (it must be "atomic")

```
int *lock = 0;

Booking()
{
    key = 1;
    while (key == 1)
        SWAP(lock, key);
    if (AvailableSeats > 0)
    {
        /* Processing for booking */
        AvailableSeats -= 1;
    }
    (*lock) = 0;
}
```

Mutual exclusion: Semaphores

```
InitSemaphore(Mutex, 1);
```

```
Booking()
```

```
{  
    P(Mutex); // Take a token
```

```
    if (AvailableSeats > 0)  
    {  
        /* Processing for booking */  
        AvailableSeats -= 1;  
    }
```

```
    V(Mutex); // Release a token  
}
```

- Semaphore
 - Distributes tokens
 - Has a maximum of tokens
 - Managed by the kernel
- Init
 - Initialize counter
- P
 - Distribute tokens
 - Block if no available token
- V
 - Take back tokens

Mutual exclusion: Semaphores

```
P(sem_t sem)
{
    disable_interrupt();

    sem.Count -= 1;
    if (sem.Count < 0)
    {
        // Block current process
        enqueue(sem.Queue, cur_PID);
        cur_PID.state = blocked;
        reschedule = true; // yield
    }

    enable_interrupt();
}
```

```
V(sem_t sem)
{
    disable_interrupt();

    sem.Count += 1;
    if (sem.Count <= 0)
    {
        // 1+ process is waiting
        process = dequeue(sem.Queue);
        process.state = running;
        reschedule = true; // yield
    }

    enable_interrupt();
}
```

Mutual exclusion: Semaphores

```
InitSemaphore(Mutex, 1);
```

```
Booking()
```

```
{
```

```
    P(Mutex); // Take a token
```

```
    if (AvailableSeats > 0)
```

```
    {
```

```
        /* Processing for booking */
```

```
        AvailableSeats -= 1;
```

```
    }
```

```
    V(Mutex); // Release a token
```

```
}
```

```
InitSemaphore(sem_t sem, int val)
```

```
{
```

```
    disable_interrupt();
```

```
    // Maximum tokens given
```

```
    sem.Count = val;
```

```
    // Waiting queue
```

```
    sem.Queue = NULL;
```

```
    enable_interrupt();
```

```
}
```

Mutual exclusion

Lot of other atomic instructions or algorithms:

- Fetch-and-Add
- Dekker's algorithm
- Peterson's algorithm
- ...

*Tips: « mutex » can be a state in which 2+ processes must be mutually excluded in order to execute a critical section...
...or it can be the surname to a simple « lock »*

Deadlocks

- 2 processes/threads try to access 2 resources that they mutually blocked
- 1st process blocked the 1st resource and... *[stopped by scheduler]*
- 2nd process blocked the 2nd resource and try to access the 1st resource *[is paused by blocking call]*
- 1st process tries to access 2nd resource *[is paused by blocking call]*

Both processes are blocked and won't run anymore

(SIGKILL is your sole friend in that case... if you can use it on these processes)

Deadlocks

Process 1

```
{  
    P(ressource1);  
    P(ressource2);  
  
    /* Usage of ressources */  
  
    V(ressource2);  
    V(ressource1);  
}
```

Process 2

```
{  
    P(ressource2);  
    P(ressource1);  
  
    /* Usage of ressources */  
  
    V(ressource1);  
    V(ressource2);  
}
```

*Extended problem to N processes:
Philosopher's diner*

Starvation

- The property assuring a process/thread to access a resource is not met
 - The process/thread trying to access a resource is « never » able to access it
- Multiple possible causes
 - Bad logic
 - Deadlock
 - No order is assured when a process is asking for a resource
(*no FIFO on the waiting queue*)
 - ...

Readers and Writers problem

- Multiple writers (write data) and readers (read data)
- Writers
 - Only 1 writer can write simultaneously
 - Nobody can read
 - Block the data (file, shared memory, ...)
- Readers
 - Multiple readers can read simultaneously
 - Nobody can write
 - 1st reader must block the data...
 - ...Last reader must wake up the writer (*if they are some waiting*)

Producers and Consumers

- 1 shared buffer and 2 indexes
 - One index for producing, one index for consuming
 - Read and write of data are sequential/in a circular FIFO
- Producers
 - Write in N cells of the buffer from the producing index
 - Cannot « produce » data if the buffer is full
 - Cannot work in the same cell of the consumers
- Consumers
 - Read N cells of the buffer from the consuming index
 - Cannot « consume » data if the buffer is empty
 - Cannot work in the same cell of the producers

Other synchronization tools

- Monitors

- A construction that uses mutex and other tools in order to be *thread-safe*
- Thread-safe: allows one resource to be safely used by multiple threads (no deadlocks can happen, neither

- Barriers

- Align multiple processes/threads on a specific point during execution
- Processes/Threads are put in a block state until the last one reaches the barrier